

VC-Dimension and Set Cover

1 Network Failure Detection

Recall that we call a set $S \subseteq V$ separable by k edges iff there exists a set $F \subseteq E$ with $|F| \leq k$ of edges such that S and \bar{S} are not connected in $(V, E \setminus F)$. In other words, the adversary could disconnect S from \bar{S} by deleting at most k edges. Let \mathcal{S}_k be the set system of all $S \subseteq V$ that are separable by k edges. We still need to show the following lemma:

Lemma 1. (V, \mathcal{S}_k) has VC-dimension at most $2k + 1$.

Proof. Because we are getting an upper bound, we need to show that no set $A \subseteq V$ with $|A| \geq 2k + 2$ is shattered by (V, \mathcal{S}_k) . That is, for every such A , there exists $A' \subseteq A$ such that $A' \neq A \cap S$ for all $S \in \mathcal{S}_k$ (i.e., there exists an $A' \subseteq A$ such that $|e_G(A, A \setminus A')| \geq k + 1$).

Lemma 2. Let G be an undirected, connected graph and $A \subseteq V$ a node set of size $|A| = 2k$. Then we can write $A = \{v_1, v_2, \dots, v_{2k-1}, v_{2k}\}$ such that for each i , there is a path P_i in G from v_{2i-1} to v_{2i} , and the P_i are edge-disjoint.

By Lemma 2, we can pair up the $2k + 2$ nodes in A and get edge-disjoint paths between the pairs. Define A' to be exactly one node from each pair (e.g., all the even-indexed nodes). To separate each pair, we must cut at least one edge from each path P_i . Because the P_i are edge-disjoint, we need to cut at least $k + 1$ edges. \square

We now prove Lemma 2:

Proof. Because more edges help us, we prove the hard case of a spanning tree of G . Then, we repeatedly delete all leaves that are not in A . If there is a pair $u, v \in A$ such that the unique $u - v$ path does not contain any degree 3 (or higher) node, pair up u, v and remove them. Otherwise, pair up leaves u, v such that the unique $u - v$ path has only one node of degree 3 (or higher) and delete them. We can show that such a pair always exists, and this produces edge-disjoint paths. \square

2 Set Cover/Hitting Set

2.1 Definitions

In the set cover problem, we are given a ground set X and a collection \mathcal{C} of subsets $S_i \subseteq X$. The goal is to find a (small) subcollection $\mathcal{D} \subseteq \mathcal{C}$ such that

$$\bigcup_{S \in \mathcal{D}} S = X.$$

In the hitting set problem, we are given a ground set X and a collection \mathcal{C} of subsets $S_i \subseteq X$. The goal is to find a (small) set $H \subseteq X$ such that

$$H \cap S \neq \emptyset \text{ for all } S \in \mathcal{C}.$$

A reduction from set cover to hitting set is to define $X' := \mathcal{C}$, $T_i := \{S \in \mathcal{C} : i \in S\}$, and $\mathcal{C}' := \{T_i : i \in X\}$. Then x' hits a set $T_i \in \mathcal{C}'$ iff $x' \in T_i$ iff $i \in S$ iff S covers i . So hitting all of \mathcal{C}' is equivalent to covering all of X .

For set cover, the greedy algorithm is an $\mathcal{O}(\log n)$ approximation and this is optimal unless $P = NP$. But maybe we get a better approximation for instances that are “less complex”. In particular, if (X, \mathcal{C}) has low VC-dimension.

Here, we will focus on constant VC-dimension (independent of n) and get an algorithm that finds a solution of cost at most $\mathcal{O}(OPT \cdot \log OPT)$.

2.2 Set Cover is still NP-hard

Could set cover be polytime solvable for constant VC-dimension? We will show that even for VC-dimension 2, set cover is still **NP**-hard because vertex cover instances also have VC-dimension 2.

In the vertex cover problem, we are given an undirected graph $G = (V, E)$. The goal is to find a (small) set $S \subseteq V$ such that each $e \in E$ has at least one endpoint in S . To prove VC-dimension, let $X = E$ and T_v be the set of edges incident of v . Then our range space is $R = \{T_v : v \in V\}$. To show that this set system has VC-dimension at most 2, we need to show that no 3-edge set $E' := \{e_1, e_2, e_3\}$ is shattered by R . We need to find $E'' \subseteq E'$ such that for all v , $E'' \neq E' \cap T_v$.

If the edges in E' don't all have a vertex in common, then there is no v with $E' = E' \cap T_v$. In other words, we cannot obtain all of E' as an intersection. If all edges in E' share a vertex v , then any set E'' with $|E''| = 2$ cannot be obtained, as $E'' = E' \cap T_u$ because the other endpoint of the edges in E'' (besides v) is different.

Lemma 3. *If a set cover instance (X, \mathcal{C}) has VC-dimension d , the corresponding hitting set instance (X', \mathcal{C}') has VC-dimension at most $2^{d+1} - 1$. If d is a constant, then this is still a constant.*

Proof. Later. □

2.3 Generalized ε -nets

Recall that an ε -net is a set N that hits all “large” sets (*i.e.*, those with probability mass at least ε). And, a hitting set is a set S that hits all sets. If we adjusted ε or the probabilities so that all sets are large, then an ε -net is exactly a hitting set.

Definition 4. To avoid some renormalization, we will slightly generalize the notion of an ε -net. Consider non-negative weights w on $\Omega = X$. We remove the assumption that $\sum_x w(x) = 1$ (so this is no longer a probability distribution). Then, N is a weighted ε -net *w.r.t.* w if

$$N \cap S \neq \emptyset \quad \forall S : w(S) \geq \varepsilon w(X).$$

Notice that if we renormalized and divided everything by $w(X)$, we would obtain the previous definition.

2.4 Initial Approach

To leverage the ε -net/hitting set insight, we assume that we know the size OPT of the smallest hitting set and also the actual hitting set H^* . Now, define the following weights:

$$w^*(x) = \begin{cases} 1 & x \in H^* \\ 0 & o.w. \end{cases}.$$

Also, let $\varepsilon = 1/OPT$. Then, the total weight $w^*(X) = OPT$ and for any set $S \in \mathcal{C}$, we have $w^*(S) = |S \cap H^*|$ and this is at least 1 because H^* is a hitting set. So,

$$w^*(S) \geq 1 = \frac{1}{OPT} \cdot OPT = \varepsilon w^*(X).$$

Therefore any weighted ε -net *w.r.t.* w^* is a hitting set.

So, if we knew w^* we could use the ε -net theorem to sample an ε -net of size

$$\Theta\left(\frac{d}{\varepsilon} \log \frac{d}{\varepsilon} + \frac{1}{\varepsilon} \log \frac{1}{\delta}\right) = \Theta(OPT \cdot \log OPT).$$

The sample would succeed with probability of at least $1 - \delta$. Therefore we have an $\mathcal{O}(\log OPT)$ approximation.

2.5 Multiplicative Weights Update

Obviously we don't know w^* or OPT . For now, continue to assume we know OPT , and we will try to "learn" w^* .

Algorithm 1 Las Vegas Multiplicative Weights Update

- 1: Set $w(x) = 1$ for all $x \in X$, and set $\varepsilon = OPT/2$.
 - 2: Find a weighted ε -net H for (X, \mathcal{C}, w) using random sampling. We can check that it is an ε -net and repeat until we get one.
 - 3: **if** H is a hitting set **then**
 - 4: Return H .
 - 5: **else**
 - 6: Let S be an arbitrary set in \mathcal{C} not hit by H .
 - 7: Double the weights of all $x \in S$ and repeat.
 - 8: **end if**
-

Theorem 5. *Algorithm 1 terminates in $\mathcal{O}(n)$ iterations and returns a hitting set of size $\mathcal{O}(OPT \cdot \log OPT)$.*

Proof. If the algorithm terminates, it returns a hitting set of size

$$\Theta\left(\frac{d}{\varepsilon} \log \frac{d}{\varepsilon} + \frac{1}{\varepsilon} \log \frac{1}{\delta}\right) = \Theta(OPT \cdot \log OPT).$$

The runtime analysis is similar to other multiplicative weights analyses. We keep track of (1) the total weight on X , and (2) the weight on H^* . We show that (2) grows faster than (1). Thus, if the algorithm runs too long, (2) will exceed (1) which is a contradiction.

(1) The total weight on X starts at $w(X) = |X| = n$. In each iteration, we double the weights on a subset S of X where $S \in \mathcal{C}$. This increases $w(X)$ by $w(S)$. Since H is an ε -net and S was not hit, we know $w(S) < \varepsilon w(X)$. After k iterations, $w(X) \leq n(1 + \varepsilon)^k$.

(2) The weight on H^* starts at $w(H^*) = |H^*| = OPT$. In each iteration, $H^* \cap S$ is nonempty because H^* is a hitting set. So at least one element of H^* has its weight doubled. To lower-bound the weight of H^* , notice that by convexity of $x \mapsto 2^x$, the weight is smallest if the lowest weight is always doubled. After k iterations, $w(H^*) \geq 2^{k/OPT} \cdot OPT$.

Clearly, after each iteration, $w(X) \geq w(H^*)$, so

$$(1 + \varepsilon)^k \cdot n \geq 2^{k/OPT} \cdot OPT.$$

Taking logs and solving for k we find

$$k \leq 4OPT \log \frac{n}{OPT} = \mathcal{O}(n).$$

□